



University
of Glasgow | School of
Computing Science

On The Scalability of ROS

Isaac Jordan

School of Computing Science
Sir Alwyn Williams Building
University of Glasgow
G12 8QQ

Level 4 Project — March 3, 2017

Abstract

Robots, distributed systems, and middleware.

Education Use Consent

I hereby give my permission for this project to be shown to other University of Glasgow students and to be distributed in an electronic format. **Please note that you are under no obligation to sign this declaration, but doing so would help future students.**

Name: _____ Signature: _____

Contents

1	Introduction	1
1.1	Context	1
1.2	Aims and Objectives	1
1.3	Achievements	2
2	Background	3
2.1	Robotics	3
2.2	Multi-Robot Systems	3
2.3	Scalability in Robotics	4
2.4	Robotic Middleware	4
2.4.1	An Overview Of Robotic Middleware	5
2.4.2	Communication	12
2.4.3	Computation	13
2.4.4	Configuration	13
2.4.5	Coordination	14
2.5	ROS (Robot Operating System)	14
2.6	Configuration of Robots	15
3	Communication Scalability	16
3.1	Scoping Experiments	16
3.1.1	Experiment 1 - Revising Existing Code	16
3.1.2	Experiment 2 - Rebooting	17
3.1.3	Experiment 3 - CPU Clock Speed	18

3.1.4	Experiment 4 - WiFi Connection	20
3.2	Realistic Data Experiments	21
3.2.1	Experiment 5 - CPU Clock Speed (Real Data)	22
3.2.2	Experiment 6 - WiFi Connection (Real Data)	25
4	Host Scalability	28
4.1	Experiment 7 - Vertical Scaling	28
4.1.1	Results	28
4.1.2	Further Investigation	30
4.2	Experiment 8 - Vertical Scaling On Car Platform	32
4.2.1	Results	32
5	Conclusion	33
5.1	Summary	33
5.2	Future Work	33
	Appendices	34
A	Experiment 2 Other Graphs	35
B	Code	37
B.1	Experiment 1 - Initial Code	37
B.1.1	Sender and Receiver	37
C	Running the Programs	38
D	Generating Random Graphs (Example Appendix)	39

Chapter 1

Introduction

1.1 Context

With the increase in availability of commodity computing hardware it is becoming more feasible to experiment with robotics at home. This increase in hobbyist roboticists is one plausible reason for the increase in open source robotics development. These community driven development projects are generally very flexible and extensible, given the large number of types of users wishing to use them in different ways: students, hobbyists, researchers, and enterprises. One such project is ROS (Robot Operating System). ROS describes itself as “a collection of tools, libraries, and conventions that aim to simplify the task of creating complex and robust robot[s]”[17]. Much of the knowledge on ROS and its performance is anecdotally acquired and shared on sites such as StackOverflow, here we present a more formal analysis of the performance and scalability of ROS both in an isolated system, and in a real robot car kit platform.

In Chapter 2 we present a walkthrough of the pertinent topics to understanding this research, as well as a review of a wide range of software packages that compete with ROS, in order to understand where ROS places in the market. We also identify the major themes seen across the variety of middlewares, broken down in to communication, computation, configuration, and coordination. Next we present an overview of ROS’s software architecture, as well as a overview of the robot car kit platform used.

In Chapter 3 we propose, execute, and evaluate the results of a series of experiments with the goal of understanding the limitations of ROS in a simple 2 host situation, with 1 lone node on each host. The first of these experiments use a very simple message format (a string), while the later ones explore how using complex data recorded from real sensors (a laser sensor, and a video camera) affects the limits of ROS communication.

With the understanding gained in the previous experiments, Chapter 4 explores how far ROS can scale on a single host (TODO: If horizontal scaling is completed, describe here), by vertically scaling the number of nodes on each host until performance degrades. This is first tested with simple Raspberry Pi’s, and then repeated on the robot car kit platform. Chapter 5 concludes.

1.2 Aims and Objectives

This project aims to provide an evaluation of ROS in a multi-robot situation, by first evaluating the simplest multi-robot case (2 robots/hosts communicating with each other, 1 node each) and looking at the limits of communication in this scenario - and attempt to identify possible causes for the limits. We then aim to vertically

scale these hosts to run many more nodes each - in order to explore the upper bound of how much communication a single ROS host can sustain (whether in terms of messages-per-second, bytes-per-second, or number of sending/receiving processes).

The intention of the research is to build a foundation of knowledge upon which further research in to communication systems of ROS can be conducted, so that future researchers need not depend on anecdotal performance estimations.

1.3 Achievements

This project makes several key contributions:

- A comprehensive review of existing robotic middlewares
-

Chapter 2

Background

2.1 Robotics

Robotics is the field concerning autonomous computer systems, usually involving physical interaction with its environment. Robotics lies at the intersection between computing science, electrical and mechanical engineering. Robotics has been common in popular culture since the 1940s[46], but due to the speed of technological growth, has quickly become widely ingrained in many industries. The first industrial use of robotics was in manufacturing, with General Motors putting the first robot into service in 1961[44]. Nowadays, with great increases in computing power, sensor accuracy, research investment, and software algorithms, example applications of robotics include manufacturing, agriculture, construction, mining, disaster recovery, medicine, health care, and surveillance[43].

Today's robotic systems generally consist of many small autonomous systems working together to form a coherent whole[31]. For example, a particular sensor (say a camera) may be constantly recording data and storing it in some buffer (erasing the oldest when full). This subsystem does not depend on any others to complete its task (recording the environment), but other subsystems may rely on its output (such as a computer vision package which needs video frame inputs). This style of robot design is becoming increasingly prevalent, with many robotic middlewares adopting this distributed approach, as discussed in Section 2.4.1.

Robotics is related to a field called cyber-physical systems (CPS) that concerns integrating computer systems with the physical world[49]. These can include Internet Of Things (IoT) networks[30] which can (for example) be used to control things around the household, such as light switches, central heating systems, and hoovering. This is distinct from robotics as CPS generally embody a 'think globally, act locally' approach[41] (such as using data from many sources to improve the CPS' performance in each), compared to robotics which utilises a more 'think locally, act locally' strategy - meaning that each instance of a robot generally makes its own decisions, and acts upon them solely.

2.2 Multi-Robot Systems

Multi-robot systems are specific instances of multi-agent systems. They represent a joint problem space between robotics, artificial intelligence, and distributed systems. Multi-robot systems as a formal concept is a recent development with a IEEE technical committee only being formed in 2014[23].

Multi-robot systems can consist of many intelligent agents (each of which may be comprised of many small autonomous systems) working to solve a task that any one system may not be able to solve alone. These multi-

robot are distinct from a multi-agent system in which individual nodes are generally stationary, as each agent in a multi-robot system is mobile[57]. Mobile robotics has been made more possible recently by advances in battery[39] and wireless communication[37] technologies. One such application of a multi-robot (and mobile) system is warehouse automation. Hamberg (2012)[45] provides a comprehensive overview of such a system, but a brief description is provided below.

In a warehouse, millions of items can be spread throughout miles of shelving and the requirement is that a random subset of items (those that have been bought) must arrive at a specific point (the delivery pick-up point) at a specific time (when the van is there). This task previously could be solved by having human pickers wander the isles searching for items - however given recent increases in the size of online shopping this is no longer feasible. Now, online retailers (such as Amazon) are using hundreds of individual robots to intelligently move the shelving around and bring the correct items to stationary human pickers[56]. This system requires the coordination of the individual robots, but they must all act independently in order to efficiently keep pace with the items ordered. However, this system still requires human pickers to lift items off the shelves, and place them in boxes. Amazon is using competitions such as the Amazon Picking Challenge to encourage researchers to develop reliable methods to automate this picking task[38], although we have still not reached the level of full automation in a production environment.

2.3 Scalability in Robotics

When creating multi-robot systems, scalability becomes an important concern[47]. At what level does a system architect choose to switch from a small number of expensive highly-powered robot systems to a larger number of simpler, cheaper robot systems. Several factors to consider include the processing power of each robot, the communication framework required to organise them, the cost of each robot, and the suitability of the problem to a multi-robot system.

Scalability can refer to many different concepts. Some define it as the ability to handle more code, others as more processes, and yet others as more hardware[32]. In this report, we will deal with the latter two definitions. We will refer to increasing the number of processes running on the same amount of hardware as ‘vertical’ scaling, and increasing the amount of hardware (e.g. number of host machines) as ‘horizontal’ scaling.

Swarm robotics is the extreme approach of creating many very simple robots which individually could not solve tasks or survive environments - however when they work together as a form of society they can work efficiently[53].

2.4 Robotic Middleware

Robot software developers want to work at a high level. Sensors are low level. (Why is middleware needed, whats available)

As mentioned previously in Section 2.1, software for robots is generally written as a collection of autonomous subsystems, often called modules. This is discussed further in Section 2.4.1.

Robotic middleware is a software infrastructure that is intended to provide convenient abstraction and communication paradigms for facilitating this multi-subsystem approach. In general, a robotics middleware would provide interfaces for defining each subsystem, and defining how each subsystem communicates with others. A specific example is the Player Project. Player provides an abstraction layer for robotic coding, which lets

developers focus more on their specific application logic, rather than boilerplate communication code[55]. The different approaches of specific middlewares is discussed later in Section 2.4.1.

A typical robot may consist of many individual sensors, for example a camera that can record 720p (a resolution of 1280x720) RGB video at 30fps (frames-per-second), a LIDAR range sensor that can measure distances of up to 30m at a frequency of 1-500Hz, a tri-axis gyroscope and accelerometer capable of measuring up to 16g of force with a measurement frequency of 40Hz. These devices use a range of sophisticated technologies to measure and analyse the physical world, each recording a different data format at a different data rate. These sensor modules are all independently designed, resulting in a wide array of different software and hardware interfaces - meaning that each robot software developer that wishes to use these sensors must create the software to communicate with these sensors, consume their data, and then write the robotic software that makes use of it. This results in a wide variety of implementations for manipulating the same data on each sensor, increasing the likelihood of implementation mistakes, misunderstandings, and wasting researchers' time.

Robotic middleware is the solution to this problem. A middleware provides a common interface design so that no matter what hardware is producing the data, the results are distributed in a consistent manner. This means that hardware manufacturers (or users) need only to implement one software system for each sensor that can then be distributed and reused amongst all users of that sensor, as long as those users are using that middleware. This means that creators of a system utilising the same middleware have easy access to the data created by a particular sensor as they know the software they create will be able to easily consume the sensor's data stream.

Another benefit of middleware is that this communication interface can be reused within the robotic system for communicating between distinct modules within the system. For example, a robot software developer may create a computer vision module that processes a video stream and returns a data stream containing a description of the objects in the video stream at each frame. When utilising a middleware, the researcher can make the result of their computer vision module available in a similar fashion to the sensor's video stream - meaning that high level software can utilise the results as if there was an 'object-detecting hardware sensor'. These levels of abstractions make software much more maintainable, and reusable.

2.4.1 An Overview Of Robotic Middleware

There is a wide variety of robotic middlewares in use currently. Many employ a free (libre) approach to software by making the source code available online, whereas others are created for commercial licensing using proprietary source. This overview predominantly covers open source middlewares, as there is a greater amount of information available on the design and implementation of these projects. Several of the covered middlewares have fallen out of use and/or development, and this has been noted where applicable. A discussion of the major aspects of the covered middlewares is presented after the table.

Name	Objective	Support	Capabilities	Supported Languages
------	-----------	---------	--------------	---------------------

ROS (Robot Operating System) [19]	<ul style="list-style-type: none"> • The goal of ROS is not to be a framework with the most features. Instead, the primary goal of ROS is to support code reuse in robotics research and development • Keep libraries ROS-agnostic • Easy to test • Scalable; appropriate for large runtime systems, and large development processes 	Large, active open source development	<ul style="list-style-type: none"> • Can be used in conjunction with other robot frameworks • Distributed framework of processes allows for executables to be individually designed, and loosely coupled at runtime • Encourages collaboration by easy package sharing • Not a realtime framework, although can work with realtime code 	Python, C++, and Lisp Experimental: Java, and Lua
MOOS (Mission Oriented Operating Suite) [9]	<ul style="list-style-type: none"> • Designed to facilitate research in the mobile robotic domain • Constitute a resilient, distributed and coordinated suite of software suitable for in-the-field deployment of sub-sea and land research robots • Process communication should be utterly robust and tolerant of the repeated stop/start of any process 	<p>Not widely used (judging by GitHub popularity)</p> <p>Development is possibly stagnating (no GitHub commits since 26th May 2016), core not updated since 11th May 2016</p>	<ul style="list-style-type: none"> • Platform independent, inter-process communication API • Sensor management • Navigation • Concurrent mission task execution • Vehicle safety management • Mission logging and replay • No P2P communication (client/server only) 	C++ (appears to have Python bindings)

KERL (Kent Erlang Robotic Library) [4]	<ul style="list-style-type: none"> • Created as a practical way of teaching Erlang. Simple API designed to let students learn Erlang, rather than learn KERL[42]. • Builds upon Player and Stage as a platform. Providing a simplified Erlang interface on the top[42]. • Contains simple single robot API for initial learning, and multi-robot APIs for advanced uses[42]. • Contains APIs for common tasks such as leader election, and broadcasting data to groups of processes[42]. 	Open Source, but not widely used. Development has been halted. No commits/updates since 2009 [5].	<ul style="list-style-type: none"> • Has provided a good starting point for robotics in Erlang[50]. • No full evaluation of suitability for production uses exists, however it is primarily an Erlang wrapper around Player (which is known to be suitable for production), thus may have a solid foundation to build upon. 	Erlang, and C
YARP (Yet Another Robot Platform) [28]	<ul style="list-style-type: none"> • Supports collection of programs communicating P2P • Extensible family of connection types (tcp, udp, multicast, local, MPI, XMLRPC,) • Flexible interfacing with hardware devices • Goal to increase the longevity of robot software projects 	Active, open source effort	<ul style="list-style-type: none"> • Data carrier method seems more flexible than ROS • General network set up seems similar to ROS. Many processes across one or more machines communicating P2P using Observer design pattern. [3] • Supports more operating systems than ROS 	SWIG (binding auto-generator)

Orocos (Open RObot Control Software) [14]	<ul style="list-style-type: none"> • Component based system design • Multi vendor (doesn't aim to solve every problem, but facilitate use of many projects) • Focus (aims to be the best free software framework for realtime control of robots and machine tools, nothing more, nothing less) 	Non-active. Still in use, but by very few people (judging by forum activity, and documentation errors (listed source host has gone down). No 'news' since 2013.	<ul style="list-style-type: none"> • Provides toolchain to create realtime robotics applications using modular, run-time configurable software components • Provides Kinematics and Dynamics Library for modelling and computation of kinematic chains, their motion specification, and interpolation (basically controlling things like robot arms) 	C++
CARMEN (Carnegie Mellon Robot Navigation Toolkit) [1]	<ul style="list-style-type: none"> • Open source collection of software for mobile robot control • Modular software to provide basic navigation functionalities, such as base and sensor control, logging, obstacle avoidance, localization, path planning, and mapping 	Discontinued, no new releases since 2008.	<ul style="list-style-type: none"> • Uses inter-process communication platform IPC • Centralised parameter server • Only supports a limited number of specific mobile robot bases. 	C and Java
Orca [13]	<ul style="list-style-type: none"> • Open source framework for developing component-based robotic systems • Goal to enable software reuse by defining commonly used interfaces • Aims to be as flexible as possible by not applying architectural constraints 	Discontinued, no new releases since 2009	<ul style="list-style-type: none"> • Provides some interfaces and implementations of commonly used components • Primarily client/server architecture, but allows for creation of distributed systems due to flexibility (not provided out of the box though) 	<p>C++, examples in Java, Python, and PHP.</p> <p>Interfaces can be compiled to C++, Java, Python, PHP, C#, Visual Basic, Ruby, and Obj C.</p>

Microsoft Robotics Developer Studio (v4) [6]	<ul style="list-style-type: none"> • Goal to make creating robotics applications very accessible • Supports visual programming (drag and drop components) • Supports simple Hello Robot to complex applications in mutli-robot scenarios 	No release since 2012	<ul style="list-style-type: none"> • RESTful (Representational state transfer) communication, services oriented runtime • Supports centralised, and decentralised communication 	C#, and Microsoft Visual Programming Language (VPL)
OpenRTM-aist [12]	<ul style="list-style-type: none"> • Open source platform to develop component oriented robotic systems. 	Seems moderately active (last release May 2016). Moderately active community (more popular in Japan)	<ul style="list-style-type: none"> • Supports communication based on Publisher/Subscriber model • Has a number of tools for robot system development 	C++, Python, Java
Player [15]	<ul style="list-style-type: none"> • Provides a clean and simple interface to the robot's sensors and actuators over the IP network 	Discontinued No commits since May 2016 No releases since 2012	<ul style="list-style-type: none"> • Supports multiple concurrent connections between devices • Supports flexible network structure (including P2P) 	Clients in C++, Tcl, Java, and Python

ROS2 [20]	<ul style="list-style-type: none"> • Target new use cases, such as multi-robot systems (providing a standard approach), embedded systems, real-time systems, non-ideal networks, and production environments [40] • Recreate ROS using existing new tech (such as Redis, WebSockets, DDS) • Overhaul of API (create consistent API without the 7+ years of backward compatibility that ROS1 has) 	Pre-release, but active daily development. Unstable but good future prospects given popularity of ROS1	<ul style="list-style-type: none"> • Improved communication resilience on poor networks utilising DDS [29] [51] • Communication overhead of DDS shown to be non-trivial for local connection. For remote, overhead is trivial but throughput depends on DDS library used [51] 	C99, C++11, Python3 Speculative: JavaScript
OpenRDK [11]	<ul style="list-style-type: none"> • Modular framework for distributed robotic systems • Communication achieved by a central ‘repository’ into which individual agents publish variables (and can store queues) • Uses URL-like addressing scheme • Focuses on mobile robots [35] 	Open source, no news since 2010, created for a single research group	<ul style="list-style-type: none"> • Created with an eye on the competition (compares it’s feature set with ORCA, OROCOS, and Player/Stage. • Has been used in multiple environments (single rescue robotic system, assistive robots) [35] • Has useful tools such as a graphical tool for remote inspection and management of modules, and also modules for logging and replaying [35] • No real-time support [10] 	C++
Miro [7]	<ul style="list-style-type: none"> • Builds upon other-widely used middlewares (ACE, TAO CORBA, Qt) to provide object-oriented abstractions [54] • Split in to 3 layers: Device, Service, and Framework • Communication achieved using CORBA client/server 	Last release was 2014	<ul style="list-style-type: none"> • Provides same capabilities as CORBA (type-safe and network-transparent interfaces) [54] • Demonstrated capabilities in multirobot environment • Does not have true OS independence (all robots used Linux), but shown that this can be ported to Solaris in 1 day [54] 	Any that have CORBA implementations

Xenomai [26]	<ul style="list-style-type: none"> • Real-time development framework (can be used to create any kind of real-time interface) • Important goals are extensibility, portability, and maintainability • Uses a dual-kernel approach to hard realtime [36] 	Sustained, active, open source development [25]	<ul style="list-style-type: none"> • Poor availability of detailed documentation and a lack of technical support [48] • Runs on top of an OS (most commonly the Linux kernel) • Shown to be suitable for 100% hard real-time applications [33] • No communication abstractions 	Preferred C [27] Possible: C++
CORBA (Common Object Request Broker Architecture) [2]	<ul style="list-style-type: none"> • Much more general than a ‘robotic middleware’, but often compared as it’s communications interface is similar to many middleware’s. • Software-based communications interface through which objects are located and accessed • OO abstractions utilising request-response in the library (via the Object Request Broker) • Uses Interface Definition Language (IDL) to define object interfaces 	Active, open source and proprietary implementations	<ul style="list-style-type: none"> • Criticised for poor implementations of the standard • Good language and OS independence • ‘Freedom from technologies’, meaning that (for example) C++ code can talk to Fortran legacy code and Java database code (and each can be changed independently without having to update the other code bases) • Strong typing of messages, reducing human error • Small overhead to adding to system (but dependent on implementation) • Has real-time implementations of related standard (realtime CORBA) 	Ada, C++, Java, COBOL, Lisp, Python, Ruby, Smalltalk Non-standard mappings exist for C#, Erlang, Perl, Tcl, Visual Basic

Urbi [24]	<ul style="list-style-type: none"> • Urbiscript aims to provide a programming experience tailored towards robotics (parallel, event-based, functional, OO, client/server, distributed) • Consists of defining modules called ‘UObject’s which are shells around regular components • These UObjects are then naturally supported by urbiscript which allows easier communication and orchestration 	Doesnt appear to be widely used, but is open source with many commits	<ul style="list-style-type: none"> • Interoperable with CORBA, RT-Middleware, openHRP (among others), thus URBI can act as a central platform to integrate other technologies • Brings many useful abstractions over other middlewares such as Player/Stage, Microsoft Robotics Studio, RT-Middleware, and CORBA. • Can move UObjects after compile-time, e.g. compile once and copy result to many hosts 	C++, Java Custom urbiscript scripting language for orchestration
-----------	---	---	--	---

Middleware designs can be broken down in to four groups of concepts: Communication, Computation, Configuration, and Coordination [34]. The majority of robotic middleware differences can be demonstrated as a difference in one of these groups. The following sections provide an overview of the approaches that the middlewares summarised above use to tackle each of these areas.

2.4.2 Communication

All modern robotic middlewares are comprised of multiple modules. In a non-trivial robotic system these distinct modules must exchange a variety of information in a complex web. These information channels usually have some desired bounds or characteristics, such as reliability (guarantees on information delivery), performance (general low latency, or some guarantees on delivery times), and overhead (is the communication significantly more expensive than building one monolithic module). The middlewares presented in the table above have used a variety of approaches to inter-component communication.

CORBA, and those middlewares built on top of CORBA utilise a remote object abstraction. This allows for inter-component communication to appear consistent whether the method caller and callee exists in the same address space, or in a remote address. The only explicit step to enable remote object communication is to share the object reference with the remote process (or component). This is achieved via an Object Request Broker (ORB), which objects can be registered with, and references retrieved from. This form of communication provides very neat code abstractions (as there is no need to modify how the object is used, only how the reference is acquired).

Other middlewares such as ROS have a more explicit communication paradigm. ROS uses a network of software nodes which can create message queues (known as a topic), which they can publish data of a specific type to [18]. Other nodes can subscribe to topics, generally registering a callback function which is called when some new data has been published to the topic. These mechanisms are described in detail in Section 2.5. This method of communication also allows for code to be identical whether or not the two communicating nodes are in the same address space or are remote, but the communication code itself is explicit. The programmer must

define when the topics are created, when data is published, and what happens in a subscriber when the data is published.

OpenRDK utilises a blackboard model. The analogy refers to the idea of many people stood around a blackboard and communicating only via writing things in different areas of the blackboard, without direct communication between the individuals. This is a form of centralised communication. In OpenRDK the blackboard is called a repository, and each module can communicate by publishing values (called properties) to the repository [10]. Properties can be simple values, or a queue object. The properties are addressed via a global hierarchical URL-like addressing scheme, similar to ROS topic addressing.

Player consists of a centralised server which connects to control clients via a standard TCP socket [16]. The client and server communicate via a set of simple messages. This is a lower-level communication paradigm which is very explicit in code, and requires careful control of shared resources as very little is provided by the library.

2.4.3 Computation

Each module created using a robotic middleware is generally tasked with some computation. That computation could be as simple as processing a value read from a sensor, and publishing it to some communication channel, or as complex as a multi-layered object classifier for images requiring large amounts of computation. Middlewares need to support these use-cases and everything in between, with a coherent, consistent infrastructure model. Middlewares generally provide a way of decoupling distinct computational tasks so that they can be individually designed, implemented, and tested - and then coupled together using the middleware's communication and coordination framework.

OpenRDK models computation inside modules [10]. Each module runs on a single thread, and multiple modules are grouped together to form an agent (a single process).

ROS has a less layered system. Computations are performed by nodes and these nodes directly communicate with each other. Each node generally performs a single focused task. A node usually consists of a single thread, but the programmer is free to spawn new threads as they require. See Section 2.5 for a more complete description of ROS.

All of the presented middlewares present a similar computation structure as the previous two examples, if they prescribe a computation structure at all. For example, CORBA says nothing about how the robotic application is structured, merely that object references must be shared via the Object Request Broker.

2.4.4 Configuration

The configuration of a robotic system using a middleware is often specified using some mechanic of the middleware. There are several distinct stages at which configuration is important, such as compile time, deployment time, and run time. Compile time configuration involves specifying what compiler settings are required, what libraries should be linked to (and their versions), and what structures and metadata should be created. At deployment time configuration involves setting up the system the robotic software is running on - such as installing libraries via a package manager, copying software to specific locations, and setting environment variables. Run time configuration has a wide variety of uses, but can consist of specifying API keys, database connection details, how many threads should be created for each process, which modules should start (and when), which components should interconnect, and how exceptions should be handled.

ROS uses XML files at compile time to resolve dependencies, export version numbers, and other miscellaneous meta information such as software license and author details. The ability to define the launch of ROS nodes is provided by the ‘roslaunch’ package. This package parses an XML file which defines which nodes should run, and also sets any required parameters on the ROS Parameter Server. OpenRDK utilises an XML configuration file to do similar tasks. It also provides some reliability functionality, such as respawning processes that have died [21]. MOOS also utilises a text file for runtime configuration called a ‘Mission file’ [52]. This mission file provides the necessary runtime parameters for components to set themselves up in a system.

ROS includes another way for new agents to configure themselves - in the form of a parameter server. This parameter server runs on the ROS master node, providing a central repository of settings. New non-master nodes can request specific parameters from this server during set up (and execution). This means that configuration files need not be modified in all running instances of the nodes, merely the information stored by the master node need be modified - providing a more centralised configuration than simple files.

2.4.5 Coordination

Most middlewares have explicit mechanisms for creating multiple concurrently running software components. These components need to exhibit some overall system behaviour, usually by working together. It is not enough for each component to independently run and share data, there needs to be some coordination to the system to provide controlled, reliable, behaviour. For this reason, some middlewares provide concurrency libraries, either explicitly, or implicitly such as by the computation and communication model or by language features.

In ROS, this concurrency is managed by running all nodes in separate threads, and designing the nodes in a reactive model. The node’s computation only occurs when there is data to process, or some task to complete.

Interconnections between components in many middlewares are achieved over TCP connections with the middleware libraries handling (de)serialization of the messages. Some middlewares such as Player are designed for a client/server architecture, with no/little intercommunication between client components, whereas ROS is designed entirely as a direct P2P network topology (with the master node mainly providing address look-up).

2.5 ROS (Robot Operating System)

The specific focal point of this investigation is ROS. ROS’s primary goal is one of sharing and collaboration. Robotic systems often use custom-created software such as driver software and higher-level algorithms like pathing. ROS creators want to make this custom created software reusable across a wide variety of platforms, reducing the amount of repeated independent development, and allowing for faster creation of useful robots. On top of this primary goal, ROS also aims to be very thin, allow libraries to be ROS-agnostic, be language independent, allow easy testing (unit and integration), and easily scale to large systems.

ROS achieves it’s primary goal via the use of packages. A ROS package contains all the information and files needed to perform one ‘task’. This can include code, datasets, and configuration files. ROS’s computation and communication units are nodes. A ROS node represents a process that performs a particular computation. A package generally contains one or more nodes. Nodes can communicate between each other directly with the use of messages, or invoke services. The ROS Master node is a particular node which must run on every ROS system. The Master provides look-up services for nodes (so that they can find each other) with a URL-like system. The Master also provides the Parameter Server. The parameter server allows for nodes to store and retrieve data at runtime from a centralised, shared dictionary.

The majority of inter-node communication is achieved using topics. A topic represents a strongly typed message bus to which one or more nodes publish messages, and zero or more nodes subscribe to receive published messages. There are no access permissions to a topic, any node can publish or subscribe as long as they use the correct data type.

Another option for inter-node communication is to use services. A service represents a restricted version of publisher/subscriber which implements a request/response interaction. When a node invokes a service, it sends a request (a message of a specific type) to the node implementing the service, and waits for the response. The response is sent back as another type (although possibly the same type).

The ROS community highly favours open source and sharing, as it aligns with the primary goal of ROS.

2.6 Configuration of Robots

The robots used in this project are 9 identical robot cars with front-wheel steering. (What is their set up)

This project involved the use of 9 identical robot cars with front-wheel steering. The cars were previously built from a Sunfounder Smart Video Car Kit for Raspberry Pi[22].

The car kit includes the physical pieces required to construct a robot car, such as a frame, gears, wheels, motors, step-down converters, and wires. The kit also includes a USB camera, and Wi-Fi adapter. It also includes a space for a Raspberry Pi (B+/2/3) to be seated. The robot cars were fitted with one Raspberry Pi 3 Model B each.

Chapter 3

Communication Scalability

The communication experiments look in to the scalability of communication in ROS. How well does it handle having to send high frequency work-loads of varying sizes.

3.1 Scoping Experiments

Scoping experiments were quick experiments using dummy data. Designed to identify useful areas to explore in more detail in realistic experiments later.

3.1.1 Experiment 1 - Revising Existing Code

The first experiment was designed to analyse the transfer time of messages between two machines at varying message frequencies. This would highlight whether there was some limit as to how often ROS could send and receive messages on it's topics.

These two machines were Raspberry Pi 3 Model Bs connected via ethernet to an Asus router.

The messages were sent and received using ROS Kinetic, running on the Raspbian OS.

The experimental setup was that one Raspberry Pi would run a ROS master node, another Raspberry Pi runs a sender program that notes down the message-sent time, and sends it to a 3rd Raspberry Pi which merely echoes the message back to the sender. Upon receiving the message, the sender/receiver notes the current time, and writes 'message X which was sent at Y, was received at Z' to a text file. The resulting Round Trip Time (RTT) for each message would therefore be the difference between the message-sent time and the message-received time.

The expected result of the experiment was that message latency (RTT) would be the same across all lower frequencies, until some bottleneck was reached that would then cause message latencies to exponentially increase due to congestion.

Code had been written prior to execute this experiment, so initially this was used B.1. However, this gave results that were contrary to the hypothesis. An increase in message frequency resulted in a reduction in message latency. A number of messages on higher frequencies were also dropped, and never received. As this was the opposite of the hypothesis, the first step was to critique the experiment code.

This review highlighted two major issues, the first was the echoing machine had a delay similar to the sender when the experiment design mandated that the echoer always respond as fast as it can. The second issue was the the maximum message queue size in ROS (how many messages can be buffered at once to compensate for a slow subscriber) was set equal to the message frequency of that run.

These issues were resolved by removing the code that executed the delay in the echoer, and by setting the maximum queue size to be equal to 1000 in every experiment (the number of messages expected to be sent).

The experiment was then repeated using this code, and the results from these runs agreed with the hypothesis.

3.1.2 Experiment 2 - Rebooting

Experiment 2 was an investigation in to whether the performance of ROS messages was affected by previous messages sent on the system. This was tested by comparing the performance of ROS messages of 5 consecutive message passing runs vs 5 message passing runs with system reboots between runs. This was repeated at 1KHz, 4KHz, 7KHz, and 10KHz frequencies for two reasons. The first was to investigate whether areas in which we observed consistent performance before would exhibit any difference between reboot and no reboot. Secondly, it would give further insight in to where the exact barrier between ‘good, consistent performance’ and ‘poor, erratic performance’ is.

Rebooting the systems involved has the opportunity to affect performance by stopping any background processes, interrupting slow processing messages from previous runs, and resetting any message caches and buffers in memory.

The result of the experiment was hypothesised to demonstrate no significant difference between rebooting and not-rebooting at any message frequency.




Figure 3.1: Experiment 2 - Mean Message Latency 1KHz Message Frequency

Figure 3.1 demonstrates that for relatively low message frequencies the mean message latency was consistently 1 - 1.5ms (the peak around message 500 in the full reboot data was due to 1 erroneous run at that data point). Figure 3.2 is characteristic of the higher frequency runs - the no reboot runs generally gave equal or better performance compared to the full reboot runs. See Appendix A for other mean graphs, and individual run graphs.

We can conclude from this that rebooting the host systems between experimental runs would result in worse performance, thus in future experiments the host systems will not be rebooted between runs.




Figure 3.2: Experiment 2 - Mean Message Latency 4KHz Message Frequency

3.1.3 Experiment 3 - CPU Clock Speed

The purpose of this experiment was to identify whether the limited CPU power of the Raspberry Pi system was playing a significant role in the bottlenecking of message communication at higher frequencies.

This is achieved by underclocking the Raspberry Pi 3 Model B CPU by 25%, and 50%. The reason to do this instead of changing the hardware (for example changing to desktop machines) is that this method keeps variables such as CPU architecture, disk speed, and software stack the same across the experiment.

The hypothesis for the experiment was that frequencies which we had previously seen high message latency for would result in even higher latency, and the maximum 'low latency' frequency would be lower as the core clock speed reduces.

One potential issue is that modern CPUs often underclock (reduce the clock speed and voltage) themselves, thus it can be hard to know exactly what CPU speed is being used throughout the experiment. Thus, several times throughout the experiment the current CPU speed was checked, and verified to be running at the specified maximum frequency.

Results

The results agreed with the hypothesis. At 100% CPU clock speed, only the highest 3 frequencies demonstrated sustained degradation of performance, however at 75% the top 5 frequencies showed, and at 50% the top 7 had reduced performance.




Figure 3.3: Experiment 3 - 100% CPU Speed, All Frequencies

Overall message latency was also increased as CPU core clock speed reduced. Even at the lowest frequency of 200Hz, 100% CPU had an average message latency of 1.188ms, 75%’s was 1.399, and 50%’s was 1.609ms. Higher message frequencies demonstrated greater differences as shown in Figure 3.4.




Figure 3.4: Experiment 3 - All CPU Speeds, All Frequencies

We can conclude from these results that the CPU performance of the host machines can be a limiting factor for high frequency ROS communication, and thus CPU utilisation is an important metric when evaluating why a ROS system is experiencing poor performance.

3.1.4 Experiment 4 - WiFi Connection

Experiment 4 investigates the effect of using a WiFi network connection for message passing experiments. The target robot car platform described in Section 2.6 utilise a WiFi connection, thus understanding the effect this will have on the communication performance is required to understand the final performance of the system.

It is also reasonable that multi-robot systems may desire each robot to be physically mobile, thus a wired network connection may not be feasible in many set-ups. It is important to understand how this technology will affect performance compared to the ideal wired situation.

The Raspberry Pi 3 Model B's utilised in these experiments contain 802.11n WiFi chips which support a 2.4GHz network connection.

The experiment consists of running the same code in both a wired (Ethernet) situation and a wireless (WiFi) situation. As before, the code will send timestamped messages from the sender host to the echoer host, which will send it back to the sender. The sender then records the message id and sent and received times to file. The code is run several times with a range of message frequencies from 200Hz to 2000Hz in 200Hz steps (200Hz, 400Hz, 600Hz, ...).

The hypothesis was that switching to WiFi would increase message latencies across all frequencies, and also reduce the maximum frequency that message latency remains consistent across the entire message stream (as in the previous experiment described in Section 3.1.3).

Results

The experimental results after 3 runs in each setup appeared to somewhat agree with the hypothesis. Overall message latency was higher across all frequencies using WiFi except 2000Hz - as shown in Figure 3.5.




Figure 3.5: Experiment 4 - Ethernet, All Frequencies

Figures 3.6 and 3.7 demonstrate averages across 3 runs for all frequencies of the message streams, for Ethernet and WiFi respectively. *It is clear from these figures that WiFi gives more erratic results than Ethernet, as*

well as overall lower performance, and thus its effects must be carefully considered when using WiFi in future experiments.




Figure 3.6: Experiment 4 - Ethernet, All Frequencies




Figure 3.7: Experiment 4 - WiFi, All Frequencies

3.2 Realistic Data Experiments

Previous experiments were presented as scoping experiments. These were designed to systematically understand which variables were of concern when running tests on ROS's communication performance. However, it was not certain that these prior conclusions would translate well to using ROS in a realistic scenario. Throughout the scoping experiments, the same message of 'hello world' was used as dummy data.

In order to verify the conclusions were correct, samples of realistic data was explored. Two likely data types were settled on, 'sensor data' and 'video data'.

Sensor data is the type of data likely to come from a physical sensor on the robot. This data is characterised

by small message sizes, such as 4KB.

Video data used was a 30Hz (30 frames-per-second) RGB video stream, with a resolution of 640 x 480 pixels. This message stream was measured to use 9.25MB/s of bandwidth, implying a message size of 308KB.

Both data sets were acquired from the MIT Stata Center dataset [8]. This dataset contains both sensor feeds (such as from a laser sensor), and video feeds (from a Kinect RGB + depth camera). These data sets are very large (20 - 50GB) which would require modification to the test system (Raspberry Pi 3s). Thus these datasets have been filtered down to 60 seconds of recording, resulting in 1791 camera images, and 1194 LaserScan readings.

3.2.1 Experiment 5 - CPU Clock Speed (Real Data)

This experiment aims to verify the results of Experiment 3 in Section 3.1.3 when using real data. This is achieved by repeating the experimental set-up and procedure, but replacing the sent message with a recorded payload.

The experiment was repeated using the previously mentioned sensor data, as well as video data.

The expectation was that sensor data (with it's relatively small message sizes) would give similar results to the dummy data used previously (a string consisting of 'hello world'), and that video data would demonstrate different performance characteristics due to the significantly different message sizes.

This is achieved using a ROS bag provided as part of the MIT Stata Center Dataset. A ROS bag is a data-structure which allows replaying of ROS topics. The topic can be replayed at the same rate it was recorded at, or any other desired rate. This allows for a variety of message frequencies to be used, as in previous experiments.

Results

The experiment appeared to confirm the results of Experiment 3. For sensor data, 100% CPU speed (1.2GHz) was the lowest latency at 7 out of 10 message frequencies, as shown in Figure 3.8, and 50% CPU speed (600MHz) was slowest at 9 out 10 frequencies.

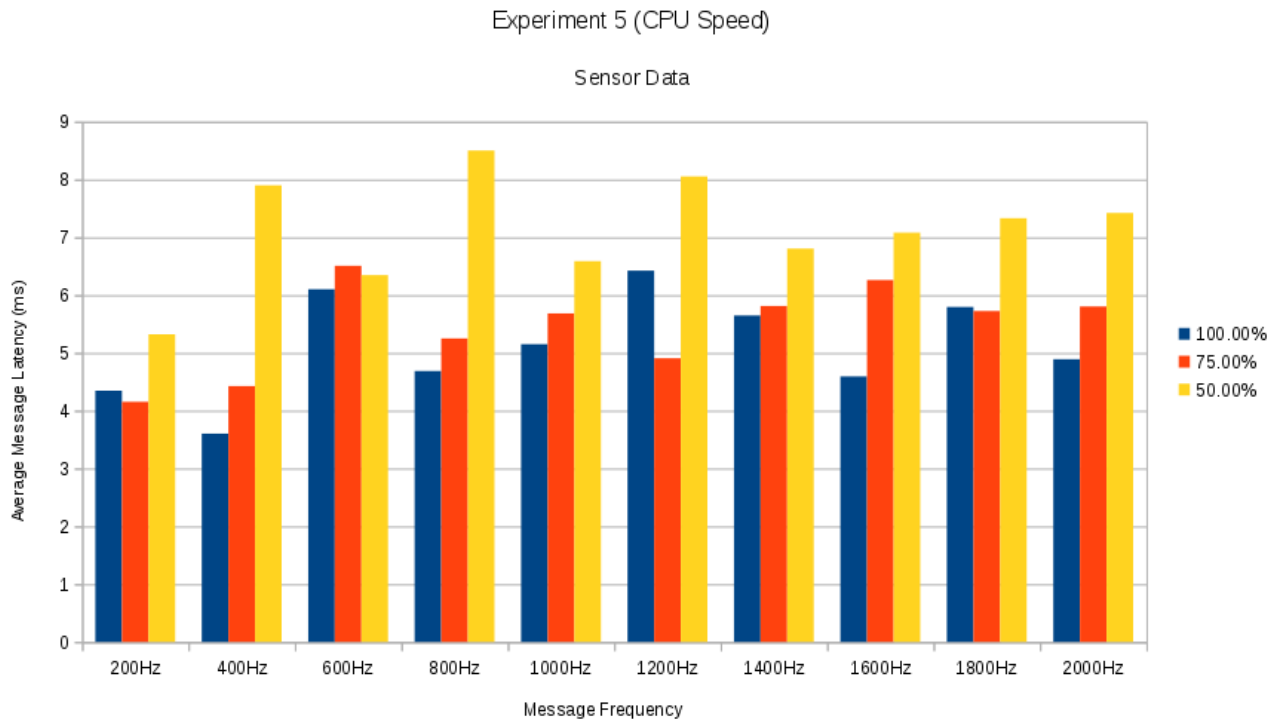


Figure 3.8: Experiment 5 - Sensor Data, All Frequencies

For video data, the trend held at the lower experimental frequencies (see Figure 3.9) with low message latencies, which stepped up slightly when CPU speed was decreased. However, the change in performance was significantly less than for the smaller message sizes of the sensor data - for example, at 20Hz the top and bottom values were only 4.9% different, whereas for sensor data at 400Hz the top and bottom values were 54% different. *This leads to the conclusion that CPU speed is less impactful as message sizes increase.* CPU speed became entirely insignificant at higher message frequencies for video data, as it is thought the network interface becomes the bottleneck. Message frequencies at larger than 30Hz gave much higher average latencies (around 4 seconds at 40Hz, compared to 60 milliseconds at 20Hz), and showed no correlation with CPU speed (see Figure 3.10).




Figure 3.9: Experiment 5 - Video Data, Low Frequencies




Figure 3.10: Experiment 5 - Video Data, High Frequencies

3.2.2 Experiment 6 - WiFi Connection (Real Data)

This experiment was a repeat of Experiment 4, except using realistic data. The data used was the same MIT State Center dataset used previously in Experiment 5.

The aim of this experiment was to verify the results of Experiment 4 were valid when using common message types. The two message types used were sensor data, and video data, as discussed previously.

Expectations from this experiment were that the results of Experiment 4 would be exacerbated. The higher latencies of wifi would be exaggerated by using larger message sizes, and performance would degrade even faster than in Experiment 4.

Results

Results showed what has been classified as ‘good’ performance at only 200Hz (see Figure 3.11). Raising the message frequency to 400Hz (and higher) introduced a significant drop in performance - from an average message latency of 23.1ms at 200Hz to 726.4ms at 400Hz, and peaking with a latency of 1721.4ms at 2000Hz (see Figure 3.12).

This leads to the conclusion that WiFi is a suitable choice of communication medium, as long as suitable low message frequency is chosen so as to ensure consistent performance. There was notable difference in performance at 200Hz from ethernet (23.1ms for WiFi, compared to 4.3ms for ethernet), however as the difference was consistent throughout the message stream (see Figure 3.11) this difference can be taken in to account for future experiments.




Figure 3.11: Experiment 6 - Sensor Data WiFi, All Frequencies




Figure 3.12: Experiment 6 - Sensor Data WiFi, All Frequencies

Further Investigation

After the results shown in Figure 3.12 demonstrated a significant difference between message frequencies of 200Hz and 400Hz, it was decided that further runs should be conducted between these frequencies to investigate the point at which performance degrades.

The extra runs conducted started at 100Hz (to gain information about performance below 200Hz), and increased in 50Hz increments up to 600Hz. The main areas of interest were 250Hz, 300Hz, and 350Hz.

As Figure 3.13 shows, message stream performance was similar as seen before. Certain streams exhibited consistent low latency throughout their streams, but above a certain frequency (350Hz and up in this case) performance begins to degrade.

Figure 3.14 more clearly shows the jump in average message latency between 300Hz and 350Hz. *This leads to the conclusion that for the sensor data message size (around 108kB), a suitable maximum frequency for publishing to ROS topics would be 300Hz.*




Figure 3.13: Experiment 6 - Sensor Data WiFi, More Frequencies




Figure 3.14: Experiment 6 - Sensor Data WiFi, More Frequencies

Chapter 4

Host Scalability

These scalability experiments look in to how well ROS scales when introduced to more nodes, or more hosts, or both.

4.1 Experiment 7 - Vertical Scaling

Prior experiments have been setting the groundwork for these scaling experiments. It is now understood how a communication intensive ROS application behaves in the simple case (2 hosts, 2 nodes), and how switching to a wireless network affects the performance.

Now we wish to understand how adding more communication nodes will affect individual message latency, and total system throughput in terms of messages sent.

This experiment will involve adding more nodes to both sender, and echoer. Each sender node will send messages to 1 echoer node, who will echo the message back to that same sender node. This allows the sender to record both sent and receive times consistently.

Increasing numbers of nodes will be added to each host, and message times recorded. The number of hosts will be a constant (2) as this experiment investigates the effect of increased ROS node counts only.

The data used will be the sensor data stream used in previous experiments.

4.1.1 Results

Figure 4.1 displays the averaged (across 3 runs) data for the 100Hz message frequency. This figure shows that performance degrades very rapidly for high node counts. At 100Hz using 8 total nodes (4 senders) results in poor performance with a message latency of 5 seconds by the end of the stream.

Figure 4.2 shows only the ‘good’ performing node counts (1 and 2 senders) for a message frequency of 100Hz. There are apparent spikes in the message latency at regular intervals. One hypothesis for these spikes is that the senders are buffering messages, perhaps to allow other system processes to use the network interface. However, these spikes do not significantly affect the wider experiment, since they are only short lived spikes.




Figure 4.1: Experiment 7 - 100Hz Message Frequency, All Node Counts




Figure 4.2: Experiment 7 - 100Hz Message Frequency, Low Node Counts

Figure 4.3 concerns a message frequency of 200Hz, and all node counts. This shows a similar pattern as Figure 4.1, with a couple notable differences.

Perhaps most interesting is the behaviour of the 256 node (128 sender) case. As expected, performance is significantly worse at the start, peaking at a message latency of 50 seconds, however after this peak the performance begins to recover - finishing at similar performance levels as the 4 node (2 sender) case. The theorised cause behind this behaviour is that at such a high node count the host system locks up and ROS nodes begin dying. In this setup, there is no supervisor for the nodes so they are not restarted. Thus, if enough nodes

die, the rest can continue functioning. This, however, is not useful behaviour in practise, as the system must halt for a long-enough time for the ROS Master's SSH connections to time out.

Another notable difference in the 200Hz case is that the 4 node (2 sender) case begins exhibiting degraded performance. This harkens back to results of Experiment 6 in Section 3.2.2, where a maximum frequency of greater than 300Hz began to show unsustainable performance.




Figure 4.3: Experiment 7 - 200Hz Message Frequency, All Node Counts

Both of these observations hold for the 300Hz message frequency scenario, with more pronounced effects in both cases. From this data, we can see the system achieves good, sustainable performance with the following set-ups: 100Hz * 1 (sending) node, 100Hz * 2 nodes, 200Hz * 1 node, and 300Hz * 1 node.

From these observations we can conclude that the results of Section 3.2.2 hold true in a more general sense, these results show that each host appears to have a communication scaling limit. This can be represented as the product of three quantites: the number of sending nodes on the host (N), the message frequency (f_m), and the message size (S_m) - giving the following equation for a Communication Scaling Limit Volume (CSLV):

$$CSLV = N \cdot f_m \cdot S_m$$

Using this equation, we can calculate the CSLV for our particular platform (Raspberry Pi 3 Model B, running ROS Kinetic) using experimental data. However, in order to demonstrate the robustness of this hypothesis, further experiments in varying the S_m (the message size) would be required.

4.1.2 Further Investigation

In order to get a clearer idea of how adding more ROS nodes affects the system, it is clear that we must use a lower message frequency which allows for sustainable performance at higher node counts.

Thus the above set-up was repeated, but with lower message frequencies. If the CSLV hypothesis were to hold true, we should be able to calculate which node counts will begin to show issues at each message frequency.

In this experiment, message frequencies of 1Hz, 10Hz, and 20Hz were used. Thus, at 1Hz we can expect greater than approximately $(300\text{Hz} / 1\text{Hz})$ 300 nodes to cause issues (we can ignore message size as we are using the same message size as previously). Respectively, we can expect the limit to lie around 30 nodes at 10Hz, and 15 nodes at 20Hz.




Figure 4.4: Experiment 8 - Logarithmic Average Message Latency, All Frequencies, All Node Counts

As Figure 4.4 shows (note the logarithmic latency scale), our expected supported node calculations were approximately correct, but somewhat underestimated the actual limits. This under-estimation is expected, as the resolution of Experiment 6 was only 50Hz (300Hz was acceptable, 350Hz had poor performance), thus the true maximum is somewhere between 300Hz and 350Hz. The figure shows that performance was acceptable at ≤ 16 nodes at 20Hz, ≤ 32 nodes at 10Hz, and ≤ 256 nodes at 1Hz.

An data series of interest is that of 256 nodes at 1Hz message frequency. We can see in Figure 4.5 the performance at this configuration is significantly worse than that at 128 nodes and less, but the performance was not as catastrophic as seen at all other frequencies (1Hz mean is 221ms, vs 10Hz mean of 41,895ms). Performance at this level was also somewhat consistent throughout the message stream, with most messages taking 200ms RTT - again, unlike other message frequencies at 256 nodes. *This indicates that a different bottleneck is possibly coming in to play. For example, perhaps the time spent switching processes contexts (between 256 different Python processes) is causing a consistent delay.* This reduced performance at 256 nodes was unexpected, according to the CSLV equation.




Figure 4.5: Experiment 8 - 1Hz Message Frequency, All Node Counts

4.2 Experiment 8 - Vertical Scaling On Car Platform

Prior experiments have been completed using what was called ‘bare’ Raspberry Pi hosts. This refers to the fact that they are off-the-shelf Raspberry Pis with no peripherals or modifications, using standard mains power. However, the test platform described in Section 2.6 represents a more realistic scenario, utilising a more unreliable power source (a battery), and several peripherals (such as a video camera). These differences in the host configuration are not expected to affect performance of the vertical scaling limits in the host (since the car platform utilises the same model of Raspberry Pi) - however, this must be experimentally confirmed.

As such, this experiment consists of a re-run of Experiment 7 (as presented in Section 4.1), but on the robot car platform. This presents some new challenges in the execution of the experiment. Due to the battery power supply the entire experiment (3 runs at 3 message frequencies with 8 different node counts) can not be run in a single battery lifetime. A single charge lasts long enough for 1.5 runs, however in order to reduce likelihood of inconsistent data it was decided to conduct the experiment 1 run at a time, with a battery charge conducted between runs - thus the experiment had to be executed over a number of days.

4.2.1 Results

Chapter 5

Conclusion

5.1 Summary

5.2 Future Work

Appendices

Appendix A

Experiment 2 Other Graphs

Other graphs commented out for now.




Figure A.1: Experiment 2 - No Reboot 1KHz Message Frequency

Appendix B

Code

B.1 Experiment 1 - Initial Code

B.1.1 Sender and Receiver

Code commented out.

Appendix C

Running the Programs

To compile this dissertation:

```
> pdflatex dissertation  
> bibtex dissertation  
> pdflatex dissertation  
  > pdflatex dissertation
```

An example of running from the command line is as follows:

```
> rosrn rosberry_experiments run_experiment.py
```

Appendix D

Generating Random Graphs (Example Appendix)

```
> java RandomGraph 100 0.9 > 100-90-00.clq
```

Bibliography

- [1] Carmen (carnegie mellon robot navigation toolkit) project homepage. <http://carmen.sourceforge.net/>. Accessed: 2017-01-27.
- [2] Corba (common object request broker architecture) project homepage. <http://www.corba.org/>. Accessed: 2017-01-27.
- [3] Getting Started with YARP Ports. http://www.yarp.it/note_ports.html. Accessed: 2017-01-27.
- [4] Kerl (kent erlang robotics library) project homepage. <http://kerl.sourceforge.net/>. Accessed: 2017-01-27.
- [5] Kerl svn repository. <https://sourceforge.net/p/kerl/code/HEAD/tree/>. Accessed: 2017-01-27.
- [6] Microsoft robotics developer studio 4 project homepage. <https://www.microsoft.com/en-gb/download/details.aspx?id=29081>. Accessed: 2017-01-27.
- [7] Miro project homepage. <http://users.isr.ist.utl.pt/~jseq/ResearchAtelier/misc/Miro%20-%20Middleware%20for%20Robots>. Accessed: 2017-01-27.
- [8] Mit stata center dataset. <https://projects.csail.mit.edu/stata/downloads.php>. Accessed: 2017-01-27.
- [9] Moos (mission oriented operating suite) project homepage. <http://www.robots.ox.ac.uk/~mobile/MOOS/wiki/pmwiki.php>. Accessed: 2017-01-27.
- [10] OpenRDk Introduction. <http://openrdk.sourceforge.net/index.php?n=Main.Introduction>. Accessed: 2017-01-27.
- [11] Openrdk project homepage. <http://openrdk.sourceforge.net/>. Accessed: 2017-01-27.
- [12] Openrtm-aist project homepage. <http://www.openrtm.org/>. Accessed: 2017-01-27.
- [13] Orca project homepage. <http://orca-robotics.sourceforge.net/>. Accessed: 2017-01-27.
- [14] Orocos (open robot control software) project homepage. <http://www.orocos.org/>. Accessed: 2017-01-27.
- [15] Player project homepage. <http://playerstage.sourceforge.net/>. Accessed: 2017-01-27.
- [16] Player server manual. <http://playerstage.sourceforge.net/doc/Player-1.6.5/player-html/>. Accessed: 2017-01-27.
- [17] Ros about page. <http://www.ros.org/about-ros/>. Accessed: 2017-03-03.
- [18] Ros concepts. <http://wiki.ros.org/ROS/Concepts>. Accessed: 2017-02-05.
- [19] Ros (robot operating system) project homepage. <http://www.ros.org/>. Accessed: 2017-01-27.

- [20] Ros2 (robot operating system) project homepage. <http://design.ros2.org/>. Accessed: 2017-01-27.
- [21] roslaunch ros package. <http://wiki.ros.org/roslaunch>. Accessed: 2017-01-27.
- [22] Sunfounder robot car kit. <https://www.sunfounder.com/rpi-car.html>. Accessed: 2017-01-27.
- [23] Technical Committee on Multi-Robot Systems (TC MRS) of the IEEE Robotics and Automation Society. <http://multirobotsystems.org/>. Accessed: 2017-01-27.
- [24] Urbi project homepage. <http://www.gostai.com/products/urbi/>. Accessed: 2017-01-27.
- [25] Xenomai Git Repositories. <https://git.xenomai.org/>. Accessed: 2017-01-27.
- [26] Xenomai project homepage. <https://xenomai.org/>. Accessed: 2017-01-27.
- [27] Xenomai tutorial. <http://www.cs.ru.nl/lab/xenomai/exercises/ex01/Exercise-1.html>. Accessed: 2017-01-27.
- [28] Yarp (yet another robot platform) project homepage. <http://yarp.it/>. Accessed: 2017-01-27.
- [29] Mateusz Macia Adam Dbrowski, Rafa Kozik. Evaluation Of ROS2 Communication Layer. <http://roscon.ros.org/2016/presentations/rafal.kozik-ros2evaluation.pdf>. Accessed: 2017-01-27.
- [30] Luigi Atzori, Antonio Iera, and Giacomo Morabito. The internet of things: A survey. *Computer networks*, 54(15):2787–2805, 2010.
- [31] B. Bauml and G. Hirzinger. Agile robot development (ard): A pragmatic approach to robotic software. In *2006 IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 3741–3748, Oct 2006.
- [32] André B. Bondi. Characteristics of scalability and their impact on performance. In *Proceedings of the 2Nd International Workshop on Software and Performance, WOSP '00*, pages 195–203, New York, NY, USA, 2000. ACM.
- [33] Jeremy H Brown and Brad Martin. How fast is fast enough? choosing between xenomai and linux for real-time applications.
- [34] Davide Brugali and Azamat Shakhimardanov. Component-based robotic engineering (part ii). *IEEE Robotics & Automation Magazine*, 17(1):100–112, 2010.
- [35] D. Calisi, A. Censi, L. Iocchi, and D. Nardi. Openrdk: A modular framework for robotic software development. In *2008 IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 1872–1877, Sept 2008.
- [36] Byoung Wook Choi, Dong Gwan Shin, Jeong Ho Park, Soo Yeong Yi, and Seet Gerald. Real-time control architecture using xenomai for intelligent service robots in usn environments. *Intelligent Service Robotics*, 2(3):139–151, 2009.
- [37] Carlos Cordeiro, Dmitry Akhmetov, and Minyoung Park. Ieee 802.11 ad: Introduction and performance evaluation of the first multi-gbps wifi technology. In *Proceedings of the 2010 ACM international workshop on mmWave communications: from circuits to networks*, pages 3–8. ACM, 2010.
- [38] Nikolaus Correll, Kostas E Bekris, Dmitry Berenson, Oliver Brock, Albert Causo, Kris Hauser, Kei Okada, Alberto Rodriguez, Joseph M Romano, and Peter R Wurman. Lessons from the amazon picking challenge. *arXiv preprint arXiv:1601.05484*, 2016.
- [39] Deborah Estrin, David Culler, Kris Pister, and Gaurav Sukhatme. Connecting the physical world with pervasive networks. *IEEE pervasive computing*, 1(1):62–63, 2002.
- [40] Brain Gerkey. Why ROS 2.0? http://design.ros2.org/articles/why_ros2.html. Accessed: 2017-01-27.

- [41] Geoff Gordon and CMU Machine Learning. Think globally, act locally.
- [42] Sten Grüner and Thomas Lorentsen. Teaching erlang using robotics and player/stage. In *Proceedings of the 8th ACM SIGPLAN workshop on ERLANG*, pages 33–40. ACM, 2009.
- [43] Martin Hägele, Klas Nilsson, J Norberto Pires, and Rainer Bischoff. Robots at work. In *Springer handbook of robotics*, pages 1385–1781. Springer, 2016.
- [44] Martin Hägele, Klas Nilsson, J Norberto Pires, and Rainer Bischoff. A short history of industrial robotics. In *Springer handbook of robotics*, pages 1386–1392. Springer, 2016.
- [45] Roelof Hamberg and Jacques Verriet. *Automation in warehouse development*. Springer, 2012.
- [46] Neil G Hockstein, CG Gourin, RA Faust, and David J Terris. A history of robots: from science fiction to surgical robotics. *Journal of robotic surgery*, 1(2):113–118, 2007.
- [47] Eric Klavins. Communication complexity of multi-robot systems. In *Algorithmic Foundations of Robotics V*, pages 275–291. Springer, 2004.
- [48] Jae Hwan Koh and Byoung Wook Choi. Real-time performance of real-time mechanisms for rtai and xenomai in various running conditions.
- [49] Edward A. Lee. Cyber physical systems: Design challenges. Technical Report UCB/EECS-2008-8, EECS Department, University of California, Berkeley, Jan 2008.
- [50] Andreea Lutac, Natalia Chechina, Gerardo Aragon-Camarasa, and Phil Trinder. Towards reliable and scalable robot communication. In *Proceedings of the 15th International Workshop on Erlang*, pages 12–23. ACM, 2016.
- [51] Yuya Maruyama, Shinpei Kato, and Takuya Azumi. Exploring the performance of ros2. In *Proceedings of the 13th International Conference on Embedded Software*, EMSOFT ’16, pages 5:1–5:10, New York, NY, USA, 2016. ACM.
- [52] Paul Newman. Introduction to programming with moos. *Communications,(November)*, pages 1–34, 2009.
- [53] Erol Şahin. Swarm robotics: From sources of inspiration to domains of application. In *International workshop on swarm robotics*, pages 10–20. Springer, 2004.
- [54] H. Utz, S. Sablatnog, S. Enderle, and G. Kraetzschmar. Miro - middleware for mobile robot applications. *IEEE Transactions on Robotics and Automation*, 18(4):493–497, Aug 2002.
- [55] Richard T Vaughan, Brian P Gerkey, and Andrew Howard. On device abstractions for portable, reusable robot code. In *Intelligent Robots and Systems, 2003.(IROS 2003). Proceedings. 2003 IEEE/RSJ International Conference on*, volume 3, pages 2421–2427. IEEE, 2003.
- [56] Peter R Wurman, Raffaello D’Andrea, and Mick Mountz. Coordinating hundreds of cooperative, autonomous vehicles in warehouses. *AI magazine*, 29(1):9, 2008.
- [57] Zhi Yan, Nicolas Jouandeau, and Arab Ali Cherif. A survey and analysis of multi-robot coordination. *International Journal of Advanced Robotic Systems*, 10(12):399, 2013.